The Claude-Mem
Festival Report Card
Judging 17 hackathon projects on how they actually used agentic memory.
data-on-a-stick
cloud sweets
August 23, 2026.
$1,000 Prize Challenge.
17 Repositories Cloned, Read, and Scored.
Gemini Notebook

# The Challenge: Build an Agent That Remembers (And the Five Judges Who Checked the Math)

JUDGE PROFILES: FIVE-POINT SCORING SYSTEM

| | |
|---|---|
| | **Depth-chan (0-5 pts):** Is Claude-Mem the architecture (5), or just a name-drop in the docs (1)? |
| | **Fit-kun (0-5 pts):** Does it hit the 7 prize categories (warm boot, retrieval, UI)? |
| | **Gear-san (0-5 pts):** Does it run? (0 = Vaporware, 3 = Runs with setup, 5 = Demoable with tests) |
| | **Coffee-sensei (0-5 pts):** The Monday Morning Tiebreaker. Would a real developer use this? |
| | **Magnifying-Glass-chan (0-5 pts):** Honesty. Does the README match the actual codebase? |

Gemini Notebook

# The Festival Leaderboard & The Hard Truths

## TL;DR Announcement Board

**No one reached Level 5 Depth.**
(Zero projects used the MCP search server, timeline, or get_observations tool).

**The plumbing leaked.**
(The most common failure was wrong port numbers or silent API failures, not lack of ambition).

**Honesty was high.**
(Most teams admitted their shortcuts!)

Gemini Notebook

# 1. Temporal (seanowww/temporal-clean)

## What it does

A provenance layer for PR review that ingests GitHub, Slack, Jira, and agent events into a graph to render an evidence-backed timeline.

## Why it wins (Yay!) ✨

- The only submission where Claude-Mem's data model is load-bearing.
- Derives the worker port perfectly.
- Honors the *hasMore* flag for pagination.
- Normalizes observations (decisions = 5 pts, changes = 2 pts).
- Falls back to reading SQLite directly with PRAGMA table_info checks to prevent schɛa drift.
- 47/47 tests pass.

## Why it loses points (Oh no...) 💧

- Claude-Mem is technically optional and not exercised in the polished demo.
- Strictly read-only.

**20/25 (First Place)**

Gemini Notebook

# 2. TeamBrain (MrFibonacc1/yc-hacks-greptileFast)

## What it does

Pools multiple engineers' local Claude-Mem databases into a shared team brain via a Bun server, injecting team briefings at session start.

## Why it wins (Yay!) ✨

- Most ambitious and category-dense.
- Speaks 3 of Claude-Mem's formats correctly.
- The only project to ship a custom mode (modes/code--team.json) using double-dash (modes/code--team.json) parent inheritance correctly.
- 198/198 tests pass.

## Why it loses points (Oh no...) 💧

- The headline import path is fatally broken on a real install.
- Prefers a legacy table (memory_items) which is empty on v13, instead of reading observations.
- Returned opsScanned: 0 on a 100k-row test database.

**18/25**
**(Second Place -**
**Conditional on a one-line table fix!)**

Gemini Notebook

# 3. RegressionForge (KaushikSiva/regression-forge)

## What it does

A post-deployment regression gate that drives Playwright through a storefront, verifies logs/webhooks, and emits a PASS/FAIL certificate.

## Why it wins (Yay!) ✨

- The only team to stand up Claude-Mem's server-beta runtime.
- Verified exact payload shapes against real route handlers.
- Real Python client hitting /v1/projects and /v1/search.
- Surfaces unavailable states honestly instead of faking them.

## Why it loses points (Oh no...) 💧

- The recall query is just a hardcoded literal string: 'RegressionForge PASS'.
- Returns its own JSON summary echoed through an index, not actual observer insights.
- Memory never touches the final deployment gate.

**16/25 (Third Place)**

Gemini Notebook

# 4. Rewind (Ashutosh-Brain/greptile-fasthackathon-26)

## What it does

A Codex-first context record tying evidence to git commits in Supabase, allowing teammates to fork corrected replays.

## Why it wins (Yay!) ✨

- Incredibly accurate client. Hits 6 different /api/ endpoints perfectly.
- Normalizes everything into a typed memory.json.
- Brutally candid README separating 'plugin installed' from 'worker connected'.

## Why it loses points (Oh no...) 💧

- It is write-only dead data.
- Nothing in the UI or backend ever reads memory.json.
- The UI shows counts, but never an actual observation.

**16/25**
**(Fourth Place - Lost on tiebreaker)**

Gemini Notebook

# 5. SourceTether (abheeshtroy/sourcetether)

## What it does

A source-verified memory gate that binds a handwritten claim to a TypeScript AST fingerprint, withholding stale memory if the code changes.

## Why it wins (Yay!) ✨

- Solves memory staleness (a massive, unaddressed field problem).
- 42 tests pass clean.
- Perfect field mapping on the GET observation call.
- Highly honest limitations section.

## Why it loses points (Oh no...) 💧

- Deliberately discards the actual observation content, keeping only the ID/Timestamp.
- Uses the hardcoded port 37701 (the author's personal UID port).

**15/25 (Fifth Place)**

Gemini Notebook

# 6. Product Git (Luciayuanzhu/product-git)

## What it does

Version control for product behavior. A web workbench lets humans accept/lock agent-declared changes, sending scoped context packs to the next turn.

## Why it wins (Yay!) ✨

- Real, tested MCP product.
- Live search call hits /api/search using package names and injects results as a 'Director Note'.

## Why it loses points (Oh no...) 💧

- Truncates search markdown at 600 characters (including table headers and pipes) instead of compressing.
- source_refs array fails because it looks for a non-existent 'results' key.
- Project filter is missing, leaking unrelated project memories into context.

**14/25 (Sixth Place)**

Gemini Notebook

# 7. Relay (thehimalayanleo/relay)

## What it does

A continuation protocol that seals JSON capsules of goals/decisions behind capability URLs to hand off investigations between human/agent pairs.

## Why it wins (Yay!) ✨

- Correct port formulas and payload field reads.
- Search call successfully extracts observation IDs from MCP-style text responses.

## Why it loses points (Oh no...) 💧

- Classic read-and-discard trap.
- Extracts the ID but throws away the memory content, replacing it with the literal string 'Claude-Mem observation N'.
- Hardcoded to only recall for projects explicitly named 'relay'.

**13/25 (Seventh Place)**

Gemini Notebook

# 8. FDE Harness (greptile-hackathon-demo/fde-harness)

## What it does

Matches Linear subtasks against 42 change recipes, scaffolds code in a git worktree, and generates customer handoff documents.

## Why it wins (Yay!) ✨

- Strong retrieval engineering.
- One real read path: directly opens SQLite to pull decision/bugfix/feature observations for a handoff markdown prompt.

## Why it loses points (Oh no...) 💦

- The prize-winning “corrections loop” reads from its own SQLite table, not Claude-Mem.
- Two of four advertised API calls simply do not exist in the code.
- Handoff script depends on private schema instead of API.

**12/25 (Eighth Place)**

Gemini Notebook

# 9. Pinocchio (akhiping/yc-greptile)

## What it does

An agent verification tool that runs deterministic detectors for test tampering and assertion weakening, emitting a contract with a “nose length” score.

## Why it wins (Yay!) ✨

- The detectors are genuinely useful.
- Upfront README admitting Claude-Mem is “an aspirational next integration surface”.

## Why it loses points (Oh no...) 💧

- The adapter cannot connect to any real install. It reads a generic “port” key instead of CLAUDE_MEM_WORKER_PORT.
- It POSTs to a read-only route.
- The live deployed demo runs in fallback mode with padded data.

**12/25 (Ninth Place)**

Gemini Notebook

# 10. Lumina (charveeeee/lumina)

## What it does

Simplifies difficult text paragraphs based on hover time. If a user struggles with a pattern 3+ times (tracked via memory), it jumps to the max rewrite.

## Why it wins (Yay!) ✨

- Perfect wiring for a 450-line app.
- Maps real validation schemas, degrades cleanly, and proves the "ingest anything" concept via non-coding telemetry.

## Why it loses points (Oh no...) 💧

- It treats Claude-Mem like a Python dictionary counter.
- Discards memory content entirely to regex-count marker strings.
- 1-second timeout means the async observer can't actually fire in time.
- Hardcoded port 37701.

**11/25 (Tenth Place)**

Gemini Notebook

# 11. Mayday (Kush614/mayday)

## What it does

A black-box flight recorder for OpenAI Codex that maps stack traces to the exact agent “false belief” that caused the error, and re-runs it in Modal.

## Why it wins (Yay!) ✨

- Incredibly polished product. 29 test cases, Playwright checks.
- Radically honest docs naming the integration as ‘pitch garnish’.

## Why it loses points (Oh no...) 💧

- Integration depth is literally 1.
- The code never touches the database, worker, or MCP tools.
- The ‘integration’ is just that the authors used Claude-Mem as their personal dev tool while writing the app.

**11/25**
**(Eleventh Place)**

Gemini Notebook

# 12. Repo Radio (nithinaru/fasthack-yc-greptile)

## What it does

A daily AI-generated Kokoro/Modal podcast reviewing fast-rising GitHub repos, featuring an on-air ‘host memory’.

## Why it wins (Yay!) ✨

- Polish! 11 baked episodes of audio.
- Genuine callback features across episodes (Episode 2 was actually about Claude-Mem itself).

## Why it loses points (Oh no...) 💧

- The worker call fails 3 different ways: Sends ?q= instead of ?query=, POSTs to a GET route, and parses a non-existent ‘snippet’ field.
- All errors are silently swallowed.
- Submits a hand-rolled memory.json but claims a ‘real working build’.

**11/25**
**(Twelfth Place)**

Gemini Notebook

# 13. Versionless (taranggoyal70/versionless)

## What it does

Clones a repo, hashes the lockfiles, lets an agent edit one file, and runs a second clone to prove the agent didn't tamper with the test suite.

## Why it wins (Yay!) ✨

- The locked-hash validation mechanism is genuinely useful and well-tested.
- Payload schemas are mapped correctly.

## Why it loses points (Oh no...) 💧

- Gates on /api/health instead of the correct /health route, rendering memory permanently unavailable.
- The live web deploy literally hardcodes fake data (provider: 'Claude-Mem', status: 'replayed'), under a banner claiming 'Repository intelligence is real'

11/25
(Thirteenth Place)

Gemini Notebook

# 14. CleanMyAgents (naturetime10/CleanMyAgents)

## What it does

A macOS menu-bar app acting as a human-in-the-loop guardian, popping up a ‘notch island’ to allow or deny an agent’s tool calls.

## Why it wins (Yay!) ✨

- Great UI design.
- A clever noise skip-list to save agent tokens.

## Why it loses points (Oh no...) 💧

- There is zero Claude-Mem code in the repository.
- The only search hits for ‘memory’ are buried inside a massive 97MB vendored clone of OpenAI’s Codex repository.
- Much of the backend is mocked or empty.

10/25
(Fourteenth Place)

Gemini Notebook

# 15. Outcome (zkortam/wingman)

## What it does

Watches deployed support agent chats to detect frustration signals, generates bounded config diffs, and self-heals by saving overrides to Supabase.

## Why it wins (Yay!) ✨

- Excellent architecture.
- The ‘prior art’ ledger is the exact right shape for a Claude-Mem search integration.

## Why it loses points (Oh no...) 💧

- The slot is empty.
- The entire integration is a 17-line class named ClaudeMemLedger that just passes through to a non-existent interface.
- It is never instantiated. A pure, empty name-drop.

10/25
(Fifteenth Place)

Gemini Notebook

# 16. LingCode CLI (Xavierhuang/Mac_cli)

## What it does

A Swift terminal CLI companion routing prompts to the LingCode macOS IDE over a Unix socket.

## Why it wins (Yay!) ✨

- Brutally honest.
- It is a legitimate, non-trivial CLI with real Swift code and makes absolutely zero false claims.

## Why it loses points (Oh no...) 💧

- Zero mentions of Claude-Mem across all 84 files.
- It uses its own unrelated .lingcode/memory.md system.
- Has no hackathon framing. Almost certainly submitted to the wrong judging pool by mistake.

7/25
(Sixteenth Place)

Gemini Notebook

# 17. Sandhāna (shreekrithi1/epicenter)

## What it does

A pre-existing proprietary SaaS for operational intelligence, complete with Stripe billing, mobile shells, and YC pitch decks.

## Why it wins (Yay!) ✨

- They built a very convincing dedicated marketing page titled ‘Sandhāna × Claude-Mem’ that claimed all 7 hackathon prize categories.

## Why it loses points (Oh no...) 💧

- Pure vaporware integration.
- Renamed existing Postgres regex-extraction endpoints to match Claude-Mem’s vocabulary.
- Inflated repo star counts.
- Claimed open-source Apache-2.0 bridges inside a strict ‘Proprietary’ repo.

7/25
(Seventeenth Place)

Gemini Notebook

# Lessons from the Academy: Where the Plumbing Broke

## Pipe Leak 1: The Port 37701 Trap

**Issue:** Hardcoding the author's local UID port instead of dynamically reading CLAUDE_MEM_WORKER_PORT.

## Pipe Leak 2: The Read-and-Discard Loop

**Issue:** Fetching observations but discarding the title, facts, and narrative, keeping only a truncated string or ID.

## The Wall of Broken Pipes

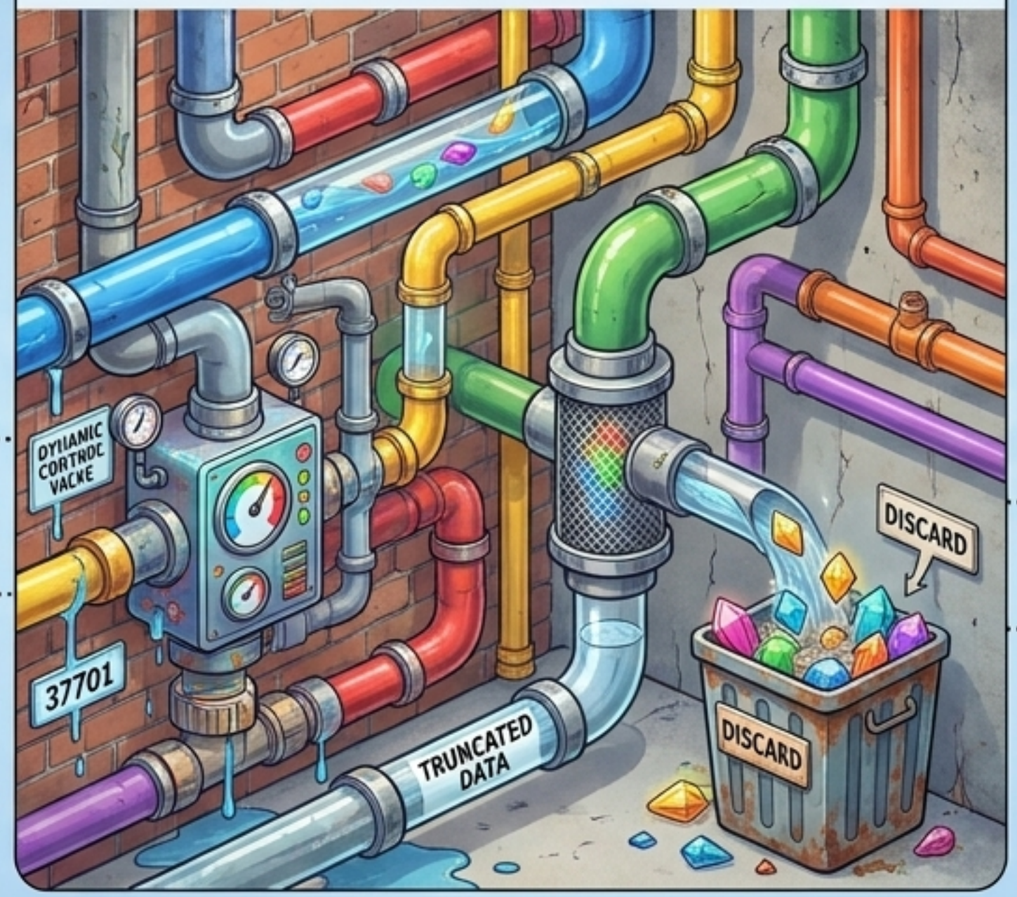


## Pipe Leak 3: The Missing Retrieval Stack

**Issue:** get_observations, mcp-search, and timeline tools were entirely unused across all 17 projects. Route typos like /api/health instead of /health killed connections.

**Building memory is hard. Wiring it is harder.** Despite the leaky pipes, the ambition on display was incredible. **Keep building, keep remembering!**

Gemini Notebook